

```

import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv("Mall_Customers.csv")

features = ['Annual Income (k$)', 'Spending Score (1-100)']
X = df[features].values

from sklearn.cluster import KMeans

wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)


C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
    warnings.warn(
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
    warnings.warn(
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
    warnings.warn(
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\

```

```
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
    warnings.warn(
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
    warnings.warn(
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
    warnings.warn(
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
```

```
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
  warnings.warn(
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
  warnings.warn(
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
  super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
  warnings.warn(

plt.plot(range(1, 11), wcss)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()
```



```
kmeans = KMeans(n_clusters=5)
y_means = kmeans.fit_predict(X)
```

```
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1412: FutureWarning: The default value of `n_init` will
change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly
to suppress the warning
```

```
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\abyji\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1436: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
    warnings.warn(
```

```
for i in range(5):
    cluster_data = df[y_means == i]
    print(f"\nCluster {i+1} characteristics:")

    print(cluster_data[features].describe())
```

```
Cluster 1 characteristics:
    Annual Income (k$)  Spending Score (1-100)
```

count	81.000000	81.000000
mean	55.296296	49.518519
std	8.988109	6.530909
min	39.000000	34.000000
25%	48.000000	44.000000
50%	54.000000	50.000000
75%	62.000000	55.000000
max	76.000000	61.000000

Cluster 2 characteristics:

	Annual Income (k\$)	Spending Score (1-100)
count	23.000000	23.000000
mean	26.304348	20.913043
std	7.893811	13.017167
min	15.000000	3.000000
25%	19.500000	9.500000
50%	25.000000	17.000000
75%	33.000000	33.500000
max	39.000000	40.000000

Cluster 3 characteristics:

	Annual Income (k\$)	Spending Score (1-100)
count	35.000000	35.000000
mean	88.200000	17.114286
std	16.399067	9.952154
min	70.000000	1.000000
25%	77.500000	10.000000
50%	85.000000	16.000000
75%	97.500000	23.500000
max	137.000000	39.000000

Cluster 4 characteristics:

	Annual Income (k\$)	Spending Score (1-100)
count	39.000000	39.000000
mean	86.538462	82.128205
std	16.312485	9.364489
min	69.000000	63.000000
25%	75.500000	74.500000
50%	79.000000	83.000000
75%	95.000000	90.000000
max	137.000000	97.000000

Cluster 5 characteristics:

	Annual Income (k\$)	Spending Score (1-100)
count	22.000000	22.000000
mean	25.727273	79.363636
std	7.566731	10.504174
min	15.000000	61.000000
25%	19.250000	73.000000
50%	24.500000	77.000000


75%	32.250000	85.750000
max	39.000000	99.000000

```
plt.scatter(X[y_means == 0,0],X[y_means ==
0,1],color='blue',label='Cluster 1')
plt.scatter(X[y_means == 1,0],X[y_means ==
1,1],color='pink',label='Cluster 2')
```

```
centroids = kmeans.cluster_centers_
```

```
plt.scatter(centroids[0, 0], centroids[0, 1], marker='*',
color='yellow', label='Centroid 1')
plt.scatter(centroids[1, 0], centroids[1, 1], marker='*',
color='yellow', label='Centroid 2')
```

```
plt.title('Clusters of customers')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.legend()
plt.show()
```



```
df["Target"]=y_means
```

```
Clustered_df=df
```

```
Clustered_df.head()
```

	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
Target					
0	1	Male	19	15	39
1					
1	2	Male	21	15	81
4					
2	3	Female	20	16	6
1					
3	4	Female	23	16	77
4					
4	5	Female	31	17	40
1					

```
X=Clustered_df.iloc[:,1:5]
```

```
y=Clustered_df.iloc[:,-1]
```

```
y.head()
```

```
0    1
1    4
2    1
3    4
4    1
```

```
Name: Target, dtype: int32
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import LabelEncoder, StandardScaler
```

```
le=LabelEncoder()
```

```
X['Gender'] = le.fit_transform(X['Gender'])
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```

```
sc=StandardScaler()
```

```
X_train=sc.fit_transform(X_train)
```

```
X_test=sc.transform(X_test)
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
gbdt = GradientBoostingClassifier(n_estimators=50,random_state=2)
```

```
gbdt.fit(X_train,y_train)
```

```
y_pred=gbdt.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
accuracy
```

0.975

```
y_pred=gbdtd.predict(X_test)
```

```
y_pred
```

```
array([0, 4, 1, 2, 2, 0, 0, 2, 2, 0, 0, 2, 3, 0, 2, 3, 0, 2, 0, 0, 2,
1,
      2, 0, 0, 0, 0, 0, 0, 3, 0, 4, 1, 0, 0, 2, 0, 3, 3, 2, 0])
```

```
prediction = pd.DataFrame({'Original Value': y_test, 'Predicted
Value': y_pred})
```

```
display(prediction)
```

	Original Value	Predicted Value
95	0	0
15	4	4
30	1	1
158	2	2
128	2	2
115	0	0
69	0	0
170	2	2
174	2	2
45	4	0
66	0	0
182	2	2
165	3	3
78	0	0
186	2	2
177	3	3
56	0	0
152	2	2
82	0	0
68	0	0
124	2	2
16	1	1
148	2	2
93	0	0
65	0	0
60	0	0
84	0	0
67	0	0
125	3	3
132	0	0
9	4	4
18	1	1
55	0	0
75	0	0
150	2	2

104	0	0
135	3	3
137	3	3
164	2	2
76	0	0